

question-1

November 15, 2023

1 ASSIGNMENT 1:

MY GITHUB LINK: [click me](#)

1.1 Question 1:

NAME: RISHAV KUMAR

ROLL NO. 2301560042

2 EX 1 : Python Pandas Exercise

```
[ ]: import pandas as pd
path =r'C:\Users\risha\Documents\KRMU\AIML_assign\datasets\Automobile_data.
      ↪csv'
df= pd.read_csv(path)
```

Exercise 1: From the given dataset print the first and last five rows

```
[ ]: df.head()
```

	index	company	body-style	wheel-base	length	engine-type \
0	0	alfa-romero	convertible	88.6	168.8	dohc
1	1	alfa-romero	convertible	88.6	168.8	dohc
2	2	alfa-romero	hatchback	94.5	171.2	ohcv
3	3	audi	sedan	99.8	176.6	ohc
4	4	audi	sedan	99.4	176.6	ohc

	num-of-cylinders	horsepower	average-mileage	price
0	four	111	21	13495.0
1	four	111	21	16500.0
2	six	154	19	16500.0
3	four	102	24	13950.0
4	five	115	18	17450.0

```
[ ]: df.tail()
```

	index	company	body-style	wheel-base	length	engine-type	\
56	81	volkswagen	sedan	97.3	171.7	ohc	
57	82	volkswagen	sedan	97.3	171.7	ohc	
58	86	volkswagen	sedan	97.3	171.7	ohc	
59	87	volvo	sedan	104.3	188.8	ohc	
60	88	volvo	wagon	104.3	188.8	ohc	

	num-of-cylinders	horsepower	average-mileage	price
56	four	85	27	7975.0
57	four	52	37	7995.0
58	four	100	26	9995.0
59	four	114	23	12940.0
60	four	114	23	13415.0

```
[ ]: path=r'C:\Users\risha\Documents\KRMU\AIML_assignment\datasets\Automobile_data.
      ↪csv'
ds= pd.read_csv(path, na_values=
      {'price': ['?', 'n.a'],
       'stroke': ['?', 'n.a'],
       'horsepower': ['?', 'n.a'],
       'peak-rpm': ['?', 'n.a'],
       'average-milage': ['?', 'n.a']})

ds.to_csv(r'C:
      ↪\Users\risha\Documents\KRMU\AIML_assignment\datasets\Automobile_data.csv')
```

Exercise 3: Find the most expensive car company name

```
[ ]: dx= df[['company', 'price']][df.price==df['price'].max()]
dx
```

	company	price
35	mercedes-benz	45400.0

```
[ ]: toyota_df= df.groupby('company')
data= toyota_df.get_group('toyota')
data
```

	index	company	body-style	wheel-base	length	engine-type	num-of-cylinders	\
48	66	toyota	hatchback	95.7	158.7	ohc	four	
49	67	toyota	hatchback	95.7	158.7	ohc	four	
50	68	toyota	hatchback	95.7	158.7	ohc	four	
51	69	toyota	wagon	95.7	169.7	ohc	four	
52	70	toyota	wagon	95.7	169.7	ohc	four	
53	71	toyota	wagon	95.7	169.7	ohc	four	
54	79	toyota	wagon	104.5	187.8	dohc	six	

	horsepower	average-mileage	price
--	------------	-----------------	-------

48	62	35	5348.0
49	62	31	6338.0
50	62	31	6488.0
51	62	31	6918.0
52	62	27	7898.0
53	62	27	8778.0
54	156	19	15750.0

Exercise 5: Count total cars per company

```
[ ]: data= df['company'].value_counts()
data
```

```
company
toyota      7
bmw         6
mazda       5
nissan       5
audi        4
mercedes-benz 4
mitsubishi  4
volkswagen  4
alfa-romero 3
chevrolet   3
honda       3
isuzu       3
jaguar      3
porsche     3
dodge       2
volvo       2
Name: count, dtype: int64
```

Exercise 6: Find each company's Highest price car

```
[ ]: categ= df.groupby('company')
data= categ['price'].max()
data
```

```
company
alfa-romero    16500.0
audi           18920.0
bmw            41315.0
chevrolet      6575.0
dodge          6377.0
honda          12945.0
isuzu          6785.0
jaguar         36000.0
mazda          18344.0
mercedes-benz  45400.0
```

```
mitsubishi      8189.0
nissan           13499.0
porsche         37028.0
toyota          15750.0
volkswagen      9995.0
volvo           13415.0
Name: price, dtype: float64
```

Exercise 7: Find the average mileage of each car making company

```
[ ]: cat_comp= df.groupby('company')
      data= cat_comp['average-mileage'].mean()
      data
```

```
company
alfa-romero      20.333333
audi             20.000000
bmw              19.000000
chevrolet        41.000000
dodge            31.000000
honda            26.333333
isuzu            33.333333
jaguar           14.333333
mazda            28.000000
mercedes-benz    18.000000
mitsubishi       29.500000
nissan            31.400000
porsche          17.000000
toyota           28.714286
volkswagen       31.750000
volvo            23.000000
Name: average-mileage, dtype: float64
```

Exercise 8: Sort all cars by Price column

```
[ ]: sort= df.sort_values(by= ['price'], ascending= False).reset_index()
      sort
```

	level_0	index	company	body-style	wheel-base	length	\
0	35	47	mercedes-benz	hardtop	112.0	199.2	
1	11	14	bmw	sedan	103.5	193.8	
2	34	46	mercedes-benz	sedan	120.9	208.1	
3	46	62	porsche	convertible	89.5	168.9	
4	12	15	bmw	sedan	110.0	197.0	
..	
56	27	36	mazda	hatchback	93.1	159.1	
57	13	16	chevrolet	hatchback	88.4	141.1	
58	22	31	isuzu	sedan	94.5	155.9	
59	23	32	isuzu	sedan	94.5	155.9	

60	47	63	porsche	hatchback	98.4	175.7
----	----	----	---------	-----------	------	-------

	engine-type	num-of-cylinders	horsepower	average-mileage	price
0	ohcv	eight	184	14	45400.0
1	ohc	six	182	16	41315.0
2	ohcv	eight	184	14	40960.0
3	ohcf	six	207	17	37028.0
4	ohc	six	182	15	36880.0
..
56	ohc	four	68	30	5195.0
57	l	three	48	47	5151.0
58	ohc	four	70	38	NaN
59	ohc	four	70	38	NaN
60	dohcv	eight	288	17	NaN

[61 rows x 11 columns]

Exercise 9: Concatenate two data frames using the following conditions

```
[ ]: germoon= {'comapny':['ford', 'mecidies', 'BMW', 'Audi'], 'price':[300, 696, 899, 454]}
japaan= {'comapny':['toyoota', 'onii', 'nii-san', 'Ohayoo'], 'price':[335, 654, 459, 934]}

df1= pd.DataFrame.from_dict(germoon)
df2= pd.DataFrame.from_dict(japaan)

main_df= pd.concat([df1, df2], keys=['Germany', 'Japan'])

main_df
```

		comapny	price
Germany	0	ford	300
	1	mecidies	696
	2	BMW	899
	3	Audi	454
Japan	0	toyoota	335
	1	onii	654
	2	nii-san	459
	3	Ohayoo	934

Exercise 10: Merge two data frames using the following condition

```
[ ]: car_price= {
    "Company": ['Toyota', 'Honda', 'BMW', 'Audi'],
    "Price" : [1234, 7653, 9874, 4982]
}
car_horsepow= {
    "Company": ['Toyota', 'Honda', 'BMW', 'Audi'],
```

```

        "horsepower": [2,4,5,8]
    }

    df1= pd.DataFrame.from_dict(car_price)
    df2= pd.DataFrame.from_dict(car_horsepow)

    data= pd.merge(df1, df2, on='Company')

    data

```

	Company	Price	horsepower
0	Toyota	1234	2
1	Honda	7653	4
2	BMW	9874	5
3	Audi	4982	8

```
[ ]:
```

3 EX 2 : Getting and Knowing your Data

This time we are going to pull data directly from the internet. Special thanks to: <https://github.com/justmarkham> for sharing the dataset and materials.

3.0.1 Step 1. Import the necessary libraries

```
[ ]: import pandas as pd
import numpy as np
```

3.0.2 Step 2. Import the dataset from this [address](#).

3.0.3 Step 3. Assign it to a variable called chipo.

```
[ ]: chipo=pd.read_csv(r'C:
↳\Users\risha\Documents\KRMU\AIML_assignment\datasets\chipotle.csv')
```

3.0.4 Step 4. See the first 10 entries

```
[ ]: chipo.head(10)
```

	order_id	quantity	item_name \
0	1	1	Chips and Fresh Tomato Salsa
1	1	1	Izze
2	1	1	Nantucket Nectar
3	1	1	Chips and Tomatillo-Green Chili Salsa
4	2	2	Chicken Bowl
5	3	1	Chicken Bowl
6	3	1	Side of Chips

7	4	1	Steak Burrito
8	4	1	Steak Soft Tacos
9	5	1	Steak Burrito

	choice_description	item_price
0	NaN	\$2.39
1	[Clementine]	\$3.39
2	[Apple]	\$3.39
3	NaN	\$2.39
4	[Tomatillo-Red Chili Salsa (Hot), [Black Beans...	\$16.98
5	[Fresh Tomato Salsa (Mild), [Rice, Cheese, Sou...	\$10.98
6	NaN	\$1.69
7	[Tomatillo Red Chili Salsa, [Fajita Vegetables...	\$11.75
8	[Tomatillo Green Chili Salsa, [Pinto Beans, Ch...	\$9.25
9	[Fresh Tomato Salsa, [Rice, Black Beans, Pinto...	\$9.25

3.0.5 Step 5. What is the number of observations in the dataset?

```
[ ]: # Solution 1
chipotle.shape[0]
```

4622

```
[ ]: # Solution 2
chipotle.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4622 entries, 0 to 4621
Data columns (total 5 columns):
#   Column                Non-Null Count  Dtype
---  -
0   order_id              4622 non-null   int64
1   quantity              4622 non-null   int64
2   item_name             4622 non-null   object
3   choice_description     3376 non-null   object
4   item_price            4622 non-null   object
dtypes: int64(2), object(3)
memory usage: 180.7+ KB
```

3.0.6 Step 6. What is the number of columns in the dataset?

```
[ ]: chipotle.shape[1]
```

5

```
[ ]: chipotle.columns
```

```
Index(['order_id', 'quantity', 'item_name', 'choice_description',
```

```
    'item_price'],
    dtype='object')
```

3.0.7 Step 8. How is the dataset indexed?

```
[ ]: chipo.index
```

```
RangeIndex(start=0, stop=4622, step=1)
```

3.0.8 Step 9. Which was the most-ordered item?

```
[ ]: c= chipo.groupby('item_name')
      c= c.sum()
      c= c.sort_values(['quantity'], ascending=False)
      c.head(1)
```

```
              order_id  quantity \
item_name
Chicken Bowl      713926         761
```

```
                                choice_description \
item_name
Chicken Bowl  [Tomatillo-Red Chili Salsa (Hot), [Black Beans...
```

```
                                item_price
item_name
Chicken Bowl  $16.98 $10.98 $11.25 $8.75 $8.49 $11.25 $8.75 ...
```

3.0.9 Step 10. For the most-ordered item, how many items were ordered?

```
[ ]: c= chipo.groupby('item_name')
      c= c.sum()
      c= c.sort_values(['quantity'], ascending=False)
      c.head(1)
```

```
              order_id  quantity \
item_name
Chicken Bowl      713926         761
```

```
                                choice_description \
item_name
Chicken Bowl  [Tomatillo-Red Chili Salsa (Hot), [Black Beans...
```

```
                                item_price
item_name
Chicken Bowl  $16.98 $10.98 $11.25 $8.75 $8.49 $11.25 $8.75 ...
```


3.0.10 Step 11. What was the most ordered item in the choice_description column?

```
[ ]: ch= chipo.groupby('choice_description').sum()
ch= ch.sort_values(['quantity'], ascending= False)
ch.head(1)
```

```
              order_id  quantity \
choice_description
[Diet Coke]          123455      159

                                     item_name \
choice_description
[Diet Coke]      Canned SodaCanned SodaCanned Soda6 Pack Soft D...

                                     item_price
choice_description
[Diet Coke]      $2.18 $1.09 $1.09 $6.49 $2.18 $1.25 $1.09 $6.4...
```

3.0.11 Step 12. How many items were ordered in total?

```
[ ]: tod= chipo.quantity.sum()
tod
```

4972

3.0.12 Step 13. Turn the item price into a float

Step 13.a. Check the item price type

```
[ ]: chipo.item_price.dtype
```

dtype('O')

Step 13.b. Create a lambda function and change the type of item price

```
[ ]: dol= lambda x:float(x[1:-1])
chipo.item_price=chipo.item_price.apply(dol)
```

Step 13.c. Check the item price type

```
[ ]: chipo.item_price.dtype
```

dtype('float64')

3.0.13 Step 14. How much was the revenue for the period in the dataset?

```
[ ]: rev= (chipo['quantity']*chipo['item_price']).sum()
print('Revenue was : '+ str(np.round(rev, 2)))
```

Revenue was : 39237.02

3.0.14 Step 15. How many orders were made in the period?

```
[ ]: order = chipo.order_id.value_counts().count()
order
```

1834

3.0.15 Step 16. What is the average revenue amount per order?

```
[ ]: # Solution 1
chipo['revenue']= chipo['quantity']*chipo['item_price']
ord= chipo.groupby(by=['order_id']).sum()
ord['revenue'].mean()
```

21.39423118865867

3.0.16 Step 17. How many different items are sold?

```
[ ]: countin= chipo.item_name.value_counts().count()
countin
```

50

```
[ ]:
```

4 EX 3 : Filtering and Sorting Data

This time we are going to pull data directly from the internet.

4.0.1 Step 1. Import the necessary libraries

```
[ ]: import pandas as pd
```

4.0.2 Step 2. Import the dataset from this [address](#).

4.0.3 Step 3. Assign it to a variable called euro12.

```
[ ]: euro12= pd.read_csv(r"C:
↳\Users\risha\Documents\KRMU\AIML_assignment\datasets\Euro_2012_stats_TEAM.
↳csv")
```

4.0.4 Step 4. Select only the Goal column.

```
[ ]: go= euro12['Goals']
go
```

```

0      4
1      4
2      4
3      5
4      3
5     10
6      5
7      6
8      2
9      2
10     6
11     1
12     5
13    12
14     5
15     2
Name: Goals, dtype: int64

```

4.0.5 Step 5. How many team participated in the Euro2012?

```
[ ]: euro12.shape[0]
```

```
16
```

4.0.6 Step 6. What is the number of columns in the dataset?

```
[ ]: euro12.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16 entries, 0 to 15
Data columns (total 35 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   Team                                16 non-null     object
 1   Goals                              16 non-null     int64
 2   Shots on target                     16 non-null     int64
 3   Shots off target                    16 non-null     int64
 4   Shooting Accuracy                   16 non-null     object
 5   % Goals-to-shots                    16 non-null     object
 6   Total shots (inc. Blocked)          16 non-null     int64
 7   Hit Woodwork                        16 non-null     int64
 8   Penalty goals                       16 non-null     int64
 9   Penalties not scored                16 non-null     int64
10   Headed goals                        16 non-null     int64
11   Passes                              16 non-null     int64
12   Passes completed                    16 non-null     int64
13   Passing Accuracy                    16 non-null     object
14   Touches                             16 non-null     int64

```

15	Crosses	16 non-null	int64
16	Dribbles	16 non-null	int64
17	Corners Taken	16 non-null	int64
18	Tackles	16 non-null	int64
19	Clearances	16 non-null	int64
20	Interceptions	16 non-null	int64
21	Clearances off line	15 non-null	float64
22	Clean Sheets	16 non-null	int64
23	Blocks	16 non-null	int64
24	Goals conceded	16 non-null	int64
25	Saves made	16 non-null	int64
26	Saves-to-shots ratio	16 non-null	object
27	Fouls Won	16 non-null	int64
28	Fouls Conceded	16 non-null	int64
29	Offsides	16 non-null	int64
30	Yellow Cards	16 non-null	int64
31	Red Cards	16 non-null	int64
32	Subs on	16 non-null	int64
33	Subs off	16 non-null	int64
34	Players Used	16 non-null	int64

dtypes: float64(1), int64(29), object(5)
memory usage: 4.5+ KB

4.0.7 Step 7. View only the columns Team, Yellow Cards and Red Cards and assign them to a dataframe called discipline

```
[ ]: dis= euro12[['Team', 'Yellow Cards', 'Red Cards']]
dis
```

	Team	Yellow Cards	Red Cards
0	Croatia	9	0
1	Czech Republic	7	0
2	Denmark	4	0
3	England	5	0
4	France	6	0
5	Germany	4	0
6	Greece	9	1
7	Italy	16	0
8	Netherlands	5	0
9	Poland	7	1
10	Portugal	12	0
11	Republic of Ireland	6	1
12	Russia	6	0
13	Spain	11	0
14	Sweden	7	0
15	Ukraine	5	0

4.0.8 Step 8. Sort the teams by Red Cards, then to Yellow Cards

```
[ ]: dis.sort_values(['Red Cards', 'Yellow Cards'], ascending=False)
dis
```

	Team	Yellow Cards	Red Cards
0	Croatia	9	0
1	Czech Republic	7	0
2	Denmark	4	0
3	England	5	0
4	France	6	0
5	Germany	4	0
6	Greece	9	1
7	Italy	16	0
8	Netherlands	5	0
9	Poland	7	1
10	Portugal	12	0
11	Republic of Ireland	6	1
12	Russia	6	0
13	Spain	11	0
14	Sweden	7	0
15	Ukraine	5	0

4.0.9 Step 9. Calculate the mean Yellow Cards given per Team

```
[ ]: round(dis['Yellow Cards'].mean())
```

7

4.0.10 Step 10. Filter teams that scored more than 6 goals

```
[ ]: euro12[euro12.Goals > 6]
```

	Team	Goals	Shots on target	Shots off target	Shooting Accuracy	\	
5	Germany	10	32	32	47.8%		
13	Spain	12	42	33	55.9%		
	% Goals-to-shots	Total shots (inc. Blocked)	Hit Woodwork	Penalty goals	\		
5	15.6%	80	2	1			
13	16.0%	100	0	1			
	Penalties not scored	...	Saves made	Saves-to-shots ratio	Fouls Won	\	
5	0	...	10	62.6%	63		
13	0	...	15	93.8%	102		
	Fouls Conceded	Offsides	Yellow Cards	Red Cards	Subs on	Subs off	\
5	49	12	4	0	15	15	
13	83	19	11	0	17	17	

	Players Used
5	17
13	18

[2 rows x 35 columns]

4.0.11 Step 11. Select the teams that start with G

```
[ ]: euro12[euro12.Team.str.startswith('G')]
```

	Team	Goals	Shots on target	Shots off target	Shooting Accuracy	\
5	Germany	10	32	32	47.8%	
6	Greece	5	8	18	30.7%	

	% Goals-to-shots	Total shots (inc. Blocked)	Hit Woodwork	Penalty goals	\
5	15.6%		80	2	1
6	19.2%		32	1	1

	Penalties not scored	...	Saves made	Saves-to-shots ratio	Fouls Won	\
5	0	...	10	62.6%	63	
6	1	...	13	65.1%	67	

	Fouls Conceded	Offsides	Yellow Cards	Red Cards	Subs on	Subs off	\
5	49	12	4	0	15	15	
6	48	12	9	1	12	12	

	Players Used
5	17
6	20

[2 rows x 35 columns]

4.0.12 Step 12. Select the first 7 columns

```
[ ]: euro12.iloc[:, 0:7]
```

	Team	Goals	Shots on target	Shots off target	\
0	Croatia	4	13	12	
1	Czech Republic	4	13	18	
2	Denmark	4	10	10	
3	England	5	11	18	
4	France	3	22	24	
5	Germany	10	32	32	
6	Greece	5	8	18	
7	Italy	6	34	45	
8	Netherlands	2	12	36	
9	Poland	2	15	23	

10	Portugal	6	22	42
11	Republic of Ireland	1	7	12
12	Russia	5	9	31
13	Spain	12	42	33
14	Sweden	5	17	19
15	Ukraine	2	7	26

	Shooting Accuracy %	Goals-to-shots	Total shots (inc. Blocked)
0	51.9%	16.0%	32
1	41.9%	12.9%	39
2	50.0%	20.0%	27
3	50.0%	17.2%	40
4	37.9%	6.5%	65
5	47.8%	15.6%	80
6	30.7%	19.2%	32
7	43.0%	7.5%	110
8	25.0%	4.1%	60
9	39.4%	5.2%	48
10	34.3%	9.3%	82
11	36.8%	5.2%	28
12	22.5%	12.5%	59
13	55.9%	16.0%	100
14	47.2%	13.8%	39
15	21.2%	6.0%	38

4.0.13 Step 13. Select all columns except the last 3.

```
[ ]: euro12.iloc[:, :-3]
```

	Team	Goals	Shots on target	Shots off target	\
0	Croatia	4	13	12	
1	Czech Republic	4	13	18	
2	Denmark	4	10	10	
3	England	5	11	18	
4	France	3	22	24	
5	Germany	10	32	32	
6	Greece	5	8	18	
7	Italy	6	34	45	
8	Netherlands	2	12	36	
9	Poland	2	15	23	
10	Portugal	6	22	42	
11	Republic of Ireland	1	7	12	
12	Russia	5	9	31	
13	Spain	12	42	33	
14	Sweden	5	17	19	
15	Ukraine	2	7	26	

Shooting Accuracy % Goals-to-shots Total shots (inc. Blocked) \

0	51.9%	16.0%	32
1	41.9%	12.9%	39
2	50.0%	20.0%	27
3	50.0%	17.2%	40
4	37.9%	6.5%	65
5	47.8%	15.6%	80
6	30.7%	19.2%	32
7	43.0%	7.5%	110
8	25.0%	4.1%	60
9	39.4%	5.2%	48
10	34.3%	9.3%	82
11	36.8%	5.2%	28
12	22.5%	12.5%	59
13	55.9%	16.0%	100
14	47.2%	13.8%	39
15	21.2%	6.0%	38

	Hit Woodwork	Penalty goals	Penalties not scored	...	Clean Sheets	\
0	0	0	0	...	0	
1	0	0	0	...	1	
2	1	0	0	...	1	
3	0	0	0	...	2	
4	1	0	0	...	1	
5	2	1	0	...	1	
6	1	1	1	...	1	
7	2	0	0	...	2	
8	2	0	0	...	0	
9	0	0	0	...	0	
10	6	0	0	...	2	
11	0	0	0	...	0	
12	2	0	0	...	0	
13	0	1	0	...	5	
14	3	0	0	...	1	
15	0	0	0	...	0	

	Blocks	Goals conceded	Saves made	Saves-to-shots ratio	Fouls Won	\
0	10	3	13	81.3%	41	
1	10	6	9	60.1%	53	
2	10	5	10	66.7%	25	
3	29	3	22	88.1%	43	
4	7	5	6	54.6%	36	
5	11	6	10	62.6%	63	
6	23	7	13	65.1%	67	
7	18	7	20	74.1%	101	
8	9	5	12	70.6%	35	
9	8	3	6	66.7%	48	
10	11	4	10	71.5%	73	
11	23	9	17	65.4%	43	

12	8	3	10	77.0%	34
13	8	1	15	93.8%	102
14	12	5	8	61.6%	35
15	4	4	13	76.5%	48

	Fouls Conceded	Offsides	Yellow Cards	Red Cards
0	62	2	9	0
1	73	8	7	0
2	38	8	4	0
3	45	6	5	0
4	51	5	6	0
5	49	12	4	0
6	48	12	9	1
7	89	16	16	0
8	30	3	5	0
9	56	3	7	1
10	90	10	12	0
11	51	11	6	1
12	43	4	6	0
13	83	19	11	0
14	51	7	7	0
15	31	4	5	0

[16 rows x 32 columns]

4.0.14 Step 14. Present only the Shooting Accuracy from England, Italy and Russia

```
[ ]: euro12.loc[euro12.Team.isin(['England', 'Italy', 'Russia']), ['Team', 'Shooting_
↳Accuracy']]
```

	Team	Shooting Accuracy
3	England	50.0%
7	Italy	43.0%
12	Russia	22.5%

```
[ ]:
```

5 EX 4 : GroupBy

5.0.1 Introduction:

GroupBy can be summarized as Split-Apply-Combine.

Special thanks to: <https://github.com/justmarkham> for sharing the dataset and materials.

Check out this [Diagram](#)

Step 1. Import the necessary libraries

```
[ ]: import pandas as pd
```

5.0.2 Step 2. Import the dataset from this [address](#).

5.0.3 Step 3. Assign it to a variable called drinks.

```
[ ]: drinks= pd.read_csv(r'C:\Users\risha\Documents\KRMU\AIML_assignment\datasets\drinks.csv')
dr=drinks
dr=dr.drop(['country'], axis=1)
```

5.0.4 Step 4. Which continent drinks more beer on average?

```
[ ]: drinks.groupby('continent').beer_servings.mean()
```

```
continent
AF      61.471698
AS      37.045455
EU     193.777778
OC      89.687500
SA     175.083333
Name: beer_servings, dtype: float64
```

5.0.5 Step 5. For each continent print the statistics for wine consumption.

```
[ ]: drinks.groupby('continent').wine_servings.describe()
```

	count	mean	std	min	25%	50%	75%	max
continent								
AF	53.0	16.264151	38.846419	0.0	1.0	2.0	13.00	233.0
AS	44.0	9.068182	21.667034	0.0	0.0	1.0	8.00	123.0
EU	45.0	142.222222	97.421738	0.0	59.0	128.0	195.00	370.0
OC	16.0	35.625000	64.555790	0.0	1.0	8.5	23.25	212.0
SA	12.0	62.416667	88.620189	1.0	3.0	12.0	98.50	221.0

5.0.6 Step 6. Print the mean alcohol consumption per continent for every column

```
[ ]: dr.groupby('continent').mean()
```

	beer_servings	spirit_servings	wine_servings
continent			
AF	61.471698	16.339623	16.264151
AS	37.045455	60.840909	9.068182
EU	193.777778	132.555556	142.222222
OC	89.687500	58.437500	35.625000
SA	175.083333	114.750000	62.416667

	total_litres_of_pure_alcohol
continent	
AF	3.007547
AS	2.170455
EU	8.617778
OC	3.381250
SA	6.308333

5.0.7 Step 7. Print the median alcohol consumption per continent for every column

```
[ ]: dr.groupby('continent').median()
```

	beer_servings	spirit_servings	wine_servings	\
continent				
AF	32.0	3.0	2.0	
AS	17.5	16.0	1.0	
EU	219.0	122.0	128.0	
OC	52.5	37.0	8.5	
SA	162.5	108.5	12.0	

	total_litres_of_pure_alcohol
continent	
AF	2.30
AS	1.20
EU	10.00
OC	1.75
SA	6.85

5.0.8 Step 8. Print the mean, min and max values for spirit consumption.

This time output a DataFrame

```
[ ]: drinks.groupby('continent').spirit_servings.agg(['mean', 'min', 'max'])
```

	mean	min	max
continent			
AF	16.339623	0	152
AS	60.840909	0	326
EU	132.555556	0	373
OC	58.437500	0	254
SA	114.750000	25	302

```
[ ]:
```

6 EX 5 : Student Alcohol Consumption

6.0.1 Introduction:

This time you will download a dataset from the UCI.

6.0.2 Step 1. Import the necessary libraries

```
[ ]: import pandas as pd
import numpy as np
```

6.0.3 Step 3. Assign it to a variable called df.

6.0.4 Step 2. Import the dataset from this [address](#).

```
[ ]: df = pd.read_csv(r'C:
↳\Users\risha\Documents\KRMU\AIML_assignment\datasets\student-mat.csv')
df.head()
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	...	\
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	...	
1	GP	F	17	U	GT3	T	1	1	at_home	other	...	
2	GP	F	15	U	LE3	T	1	1	at_home	other	...	
3	GP	F	15	U	GT3	T	4	2	health	services	...	
4	GP	F	16	U	GT3	T	3	3	other	other	...	

	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3
0	4	3	4	1	1	3	6	5	6	6
1	5	3	3	1	1	3	4	5	5	6
2	4	3	2	2	3	3	10	7	8	10
3	3	2	2	1	1	5	2	15	14	15
4	4	3	2	1	2	5	4	6	10	10

[5 rows x 33 columns]

6.0.5 Step 4. For the purpose of this exercise slice the dataframe from 'school' until the 'guardian' column

```
[ ]: st = df.loc[:, "school":"guardian"]
st.head()
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	\
0	GP	F	18	U	GT3	A	4	4	at_home	teacher	
1	GP	F	17	U	GT3	T	1	1	at_home	other	
2	GP	F	15	U	LE3	T	1	1	at_home	other	
3	GP	F	15	U	GT3	T	4	2	health	services	
4	GP	F	16	U	GT3	T	3	3	other	other	

	reason	guardian
0	course	mother
1	course	father
2	other	mother
3	home	mother
4	home	father

6.0.6 Step 5. Create a lambda function that will capitalize strings.

```
[ ]: cap = lambda x: x.capitalize()
```

6.0.7 Step 6. Capitalize both Mjob and Fjob

```
[ ]: st['Mjob'].apply(cap)
     st['Fjob'].apply(cap)
```

```
0      Teacher
1         Other
2         Other
3     Services
4         Other
...
390    Services
391    Services
392         Other
393         Other
394     At_home
Name: Fjob, Length: 395, dtype: object
```

6.0.8 Step 7. Print the last elements of the data set.

```
[ ]: st.tail()
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	\
390	MS	M	20	U	LE3	A	2	2	services	services	
391	MS	M	17	U	LE3	T	3	1	services	services	
392	MS	M	21	R	GT3	T	1	1	other	other	
393	MS	M	18	R	LE3	T	3	2	services	other	
394	MS	M	19	U	LE3	T	1	1	other	at_home	

	reason	guardian
390	course	other
391	course	mother
392	course	other
393	course	mother
394	course	father

6.0.9 Step 8. Did you notice the original dataframe is still lowercase? Why is that? Fix it and capitalize Mjob and Fjob.

```
[ ]: st['Mjob'] = st['Mjob'].apply(cap)
     st['Fjob'] = st['Fjob'].apply(cap)
     st.tail()
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	\
390	MS	M	20	U	LE3	A	2	2	Services	Services	
391	MS	M	17	U	LE3	T	3	1	Services	Services	
392	MS	M	21	R	GT3	T	1	1	Other	Other	
393	MS	M	18	R	LE3	T	3	2	Services	Other	
394	MS	M	19	U	LE3	T	1	1	Other	At_home	

	reason	guardian
390	course	other
391	course	mother
392	course	other
393	course	mother
394	course	father

6.0.10 Step 9. Create a function called `majority` that returns a boolean value to a new column called `legal_drinker` (Consider majority as older than 17 years old)

```
[ ]: def majority(x):
      if x > 17:
          return True
      else:
          return False

[ ]: st['legal_drinker'] = st['age'].apply(majority)
st.head()
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	\
0	GP	F	18	U	GT3	A	4	4	At_home	Teacher	
1	GP	F	17	U	GT3	T	1	1	At_home	Other	
2	GP	F	15	U	LE3	T	1	1	At_home	Other	
3	GP	F	15	U	GT3	T	4	2	Health	Services	
4	GP	F	16	U	GT3	T	3	3	Other	Other	

	reason	guardian	legal_drinker
0	course	mother	True
1	course	father	False
2	other	mother	False
3	home	mother	False
4	home	father	False

6.0.11 Step 10. Multiply every number of the dataset by 10.

I know this makes no sense, don't forget it is just an exercise

```
[ ]: def times10(x):
      if type(x) is int:
          return 10 * x
```

```
return x
```

```
[ ]: st.map(times10).head(10)
```

	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	\
0	GP	F	180	U	GT3	A	40	40	At_home	Teacher	
1	GP	F	170	U	GT3	T	10	10	At_home	Other	
2	GP	F	150	U	LE3	T	10	10	At_home	Other	
3	GP	F	150	U	GT3	T	40	20	Health	Services	
4	GP	F	160	U	GT3	T	30	30	Other	Other	
5	GP	M	160	U	LE3	T	40	30	Services	Other	
6	GP	M	160	U	LE3	T	20	20	Other	Other	
7	GP	F	170	U	GT3	A	40	40	Other	Teacher	
8	GP	M	150	U	LE3	A	30	20	Services	Other	
9	GP	M	150	U	GT3	T	30	40	Other	Other	

	reason	guardian	legal_drinker
0	course	mother	True
1	course	father	False
2	other	mother	False
3	home	mother	False
4	home	father	False
5	reputation	mother	False
6	home	mother	False
7	home	mother	False
8	home	mother	False
9	home	mother	False

```
[ ]:
```

7 EX 6 : MPG Cars

7.0.1 Introduction:

The following exercise utilizes data from [UC Irvine Machine Learning Repository](#)

7.0.2 Step 1. Import the necessary libraries

```
[ ]: import pandas as pd
import numpy as np
```

7.0.3 Step 2. Import the first dataset `cars1` and `cars2`.

Step 3. Assign each to a variable called `cars1` and `cars2`

```
[ ]:
```

```
cars1= pd.read_csv(r'C:
↳\Users\risha\Documents\KRMU\AIML_assignment\datasets\cars1.csv')
cars2= pd.read_csv(r'C:
↳\Users\risha\Documents\KRMU\AIML_assignment\datasets\cars2.csv')
```

7.0.4 Step 4. Oops, it seems our first dataset has some unnamed blank columns, fix cars1

```
[ ]: cars1 = cars1.loc[:, "mpg":"car"]
cars1.head()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model	\
0	18.0	8	307	130	3504	12.0	70	
1	15.0	8	350	165	3693	11.5	70	
2	18.0	8	318	150	3436	11.0	70	
3	16.0	8	304	150	3433	12.0	70	
4	17.0	8	302	140	3449	10.5	70	

	origin	car
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

7.0.5 Step 5. What is the number of observations in each dataset?

```
[ ]: print(cars1.shape)
print(cars2.shape)
```

(198, 9)

(200, 9)

7.0.6 Step 6. Join cars1 and cars2 into a single DataFrame called cars

```
[ ]: cars= pd.concat([cars1, cars2])
cars
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model	\
0	18.0	8	307	130	3504	12.0	70	
1	15.0	8	350	165	3693	11.5	70	
2	18.0	8	318	150	3436	11.0	70	
3	16.0	8	304	150	3433	12.0	70	
4	17.0	8	302	140	3449	10.5	70	
..	
195	27.0	4	140	86	2790	15.6	82	
196	44.0	4	97	52	2130	24.6	82	

197	32.0	4	135	84	2295	11.6	82
198	28.0	4	120	79	2625	18.6	82
199	31.0	4	119	82	2720	19.4	82

	origin	car
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino
..
195	1	ford mustang gl
196	2	vw pickup
197	1	dodge rampage
198	1	ford ranger
199	1	chevy s-10

[398 rows x 9 columns]

7.0.7 Step 7. Oops, there is a column missing, called owners. Create a random number Series from 15,000 to 73,000.

```
[ ]: nr_owners = np.random.randint(15000, high=73001, size=398, dtype='i')
nr_owners
```

```
array([65936, 60313, 64074, 19727, 46458, 32859, 15854, 25202, 61591,
       71275, 69046, 42269, 43071, 71470, 33756, 64702, 67920, 28883,
       42901, 37801, 32844, 31589, 25355, 37036, 62024, 40424, 42214,
       51227, 67386, 35957, 68596, 62853, 59400, 58149, 15766, 20884,
       38316, 31332, 49183, 52629, 58666, 27067, 42665, 21323, 25310,
       16915, 27009, 34352, 27182, 38928, 35607, 42473, 64360, 35536,
       63309, 39607, 17950, 21137, 24150, 18726, 38823, 42662, 65610,
       37072, 45130, 15826, 44513, 30955, 30146, 60154, 28394, 69465,
       58242, 22564, 43992, 19830, 54167, 21301, 64035, 38726, 61014,
       58021, 26688, 28023, 28208, 21168, 16239, 45498, 55714, 40814,
       57083, 48080, 23943, 40653, 32532, 68711, 60197, 59249, 24013,
       49195, 60156, 69382, 50582, 25031, 35913, 57778, 33459, 37708,
       21392, 50280, 17308, 35473, 49947, 67387, 43350, 67936, 24651,
       32968, 48698, 24003, 64259, 24320, 25793, 44880, 45540, 44127,
       55030, 63775, 36094, 35085, 32179, 31563, 44832, 42522, 68647,
       58826, 16599, 39432, 29608, 50629, 61549, 52827, 40926, 34532,
       71371, 64723, 47175, 30128, 54753, 40464, 47399, 42144, 62229,
       27922, 29076, 34164, 51387, 18319, 40510, 58262, 53211, 44960,
       20022, 49345, 49929, 53941, 66550, 66695, 19150, 71361, 45789,
       27849, 51603, 35294, 61627, 30242, 34935, 24233, 60856, 34499,
       38347, 27096, 58580, 42339, 20847, 72874, 36260, 29927, 25658,
       32956, 26488, 18581, 49463, 33759, 39963, 58050, 41653, 21919,
       51689, 35537, 34726, 55749, 64014, 27145, 65419, 57077, 65605,
```

```

15594, 41119, 33782, 55997, 69149, 50644, 43761, 30912, 57679,
20446, 34101, 22717, 63875, 70576, 35875, 39259, 25748, 15522,
52394, 23511, 42116, 25723, 30822, 26037, 21048, 68679, 31401,
38908, 66014, 17188, 41575, 52715, 53340, 37054, 43863, 63587,
60334, 48631, 50993, 51280, 31021, 25787, 61185, 40604, 53679,
31365, 69889, 48046, 55310, 32009, 31170, 17010, 59147, 18216,
51884, 61649, 36653, 71332, 16496, 26149, 15042, 15894, 33214,
46400, 27251, 63621, 60874, 30420, 52260, 70297, 50180, 40689,
27168, 48477, 19023, 32963, 69731, 52837, 56693, 39096, 41045,
42696, 62822, 38116, 18595, 46404, 16834, 52438, 67402, 63948,
56535, 71791, 44752, 46813, 54498, 53262, 37212, 57063, 56622,
53087, 66948, 25728, 18137, 72221, 25551, 29372, 69172, 67487,
50871, 68969, 18025, 61123, 40303, 70533, 33191, 18511, 55802,
18326, 59010, 53900, 20022, 62234, 67857, 49132, 59838, 28700,
19316, 34012, 29589, 57902, 55858, 64641, 43533, 33249, 49912,
50643, 52258, 26276, 52595, 59682, 35905, 58119, 37653, 35764,
60345, 36076, 61023, 31121, 36629, 30196, 60624, 65721, 42995,
33627, 54700, 59084, 15029, 70851, 35956, 60479, 62738, 46377,
46942, 23015, 59618, 33611, 60746, 56107, 35091, 57988, 64651,
25399, 22613, 70639, 33401, 59906, 16954, 30503, 16979, 55529,
70457, 49810, 48253, 18524, 53372, 62029, 45074, 17187, 21994,
48616, 46251])

```

7.0.8 Step 8. Add the column owners to cars

```
[ ]: cars['owners'] = nr_owners
cars.tail()
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model	\
195	27.0	4	140	86	2790	15.6	82	
196	44.0	4	97	52	2130	24.6	82	
197	32.0	4	135	84	2295	11.6	82	
198	28.0	4	120	79	2625	18.6	82	
199	31.0	4	119	82	2720	19.4	82	

	origin	car	owners
195	1	ford mustang gl	45074
196	2	vw pickup	17187
197	1	dodge rampage	21994
198	1	ford ranger	48616
199	1	chevy s-10	46251

8 EX 7 : Online Retails Purchase

8.0.1 Introduction:

8.0.2 Step 1. Import the necessary libraries

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

%matplotlib inline
sns.set(style="ticks")
```

8.0.3 Step 2. Import the dataset from this [address](#).


8.0.4 Step 3. Assign it to a variable called online_rt

Note: if you receive a utf-8 decode error, set `encoding = 'latin1'` in `pd.read_csv()`.

```
[ ]: path=r'C:\Users\risha\Documents\KRMU\AIML_assignment\datasets\online_Retail.csv'
online_rt=pd.read_csv(path, encoding = 'latin1')
```

8.0.5 Step 4. Create a histogram with the 10 countries that have the most 'Quantity' ordered except UK

```
[ ]: countries = online_rt.groupby(["Country"]).sum()
countries= countries.sort_values(by= 'Quantity', ascending= False)[1:11]
countries['Quantity'].plot(kind='bar')
plt.xlabel('Countries')
plt.ylabel('Quantity')
plt.title("top 10 countries with most orders")
plt.show()
```



8.0.6 Step 5. Exclude negative Quantity entries

```
[ ]: online_rt= online_rt[online_rt.Quantity>0]
      online_rt.head()
```

	InvoiceNo	StockCode	Description	Quantity	\
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	
1	536365	71053	WHITE METAL LANTERN	6	
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	

	InvoiceDate	UnitPrice	CustomerID	Country
0	12/1/10 8:26	2.55	17850.0	United Kingdom
1	12/1/10 8:26	3.39	17850.0	United Kingdom
2	12/1/10 8:26	2.75	17850.0	United Kingdom
3	12/1/10 8:26	3.39	17850.0	United Kingdom

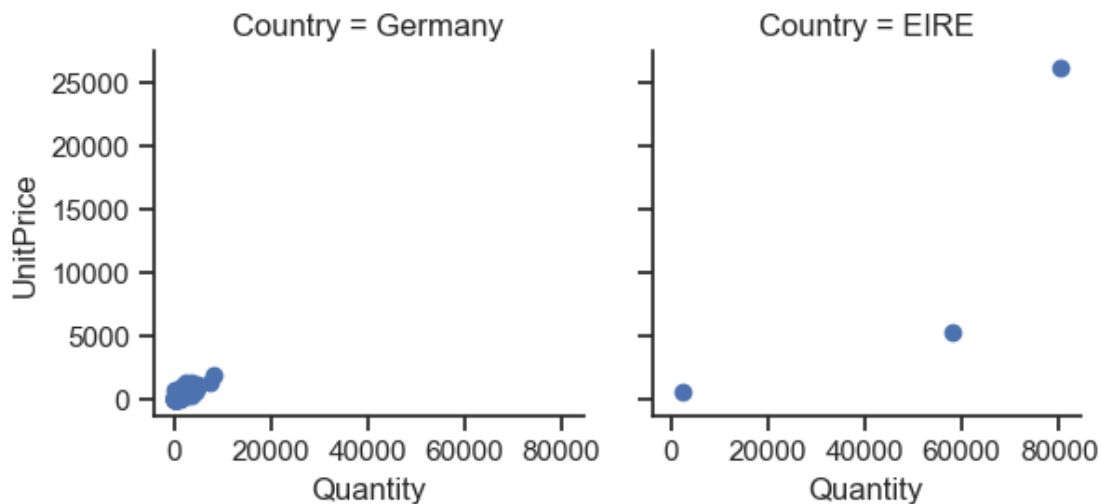
4 12/1/10 8:26 3.39 17850.0 United Kingdom

8.0.7 Step 6. Create a scatterplot with the Quantity per UnitPrice by CustomerID for the top 3 Countries (except UK)

```
[ ]: customers= online_rt.groupby(['CustomerID', 'Country']).sum()
customers= customers[customers.UnitPrice>0]
customers['Country']= customers.index.get_level_values(1)
top_countries= ['Netherland', 'EIRE', 'Germany']
customers= customers[customers['Country'].isin(top_countries)]

gs= sns.FacetGrid(customers, col='Country')
gs.map(plt.scatter, 'Quantity', 'UnitPrice', alpha=1)
gs.add_legend()
```

<seaborn.axisgrid.FacetGrid at 0x1af130ecce0>



8.0.8 Step 7. Investigate why the previous results look so uninformative.

This section might seem a bit tedious to go through. But I've thought of it as some kind of a simulation of problems one might encounter when dealing with data and other people. Besides there is a prize at the end (i.e. Section 8).

(But feel free to jump right ahead into Section 8 if you want; it doesn't require that you finish this section.)

Step 7.1 Look at the first line of code in Step 6. And try to figure out if it leads to any kind of problem.

Step 7.1.1 Display the first few rows of that DataFrame.

```
[ ]: customers= online_rt.groupby(['CustomerID', 'Country']).sum().head()
customers
```

CustomerID	Country	InvoiceNo \
12346.0	United Kingdom	541431
12347.0	Iceland	5376265376265376265376265376265376265376265376...
12348.0	Finland	5393185393185393185393185393185393185393185393...
12349.0	Italy	5776095776095776095776095776095776095776095776...
12350.0	Norway	5430375430375430375430375430375430375430375430...

CustomerID	Country	StockCode \
12346.0	United Kingdom	23166
12347.0	Iceland	8511622375714772249222771227722277322774227752...
12348.0	Finland	8499222951849918499121213212132261621981219822...
12349.0	Italy	2311223460215642141121563221312219548194849782...
12350.0	Norway	219082241279066K79191C2234884086C2255122557218...

CustomerID	Country	Description \
12346.0	United Kingdom	MEDIUM CERAMIC TOP STORAGE JAR
12347.0	Iceland	BLACK CANDELABRA T-LIGHT HOLDERAIRLINE BAG VIN...
12348.0	Finland	72 SWEETHEART FAIRY CAKE CASES60 CAKE CASES DO...
12349.0	Italy	PARISIENNE CURIO CABINETSWEEETHEART WALL TIDY P...
12350.0	Norway	CHOCOLATE THIS WAY METAL SIGNMETAL SIGN NEIGHB...

CustomerID	Country	Quantity \
12346.0	United Kingdom	74215
12347.0	Iceland	2458
12348.0	Finland	2341
12349.0	Italy	631
12350.0	Norway	197

CustomerID	Country	InvoiceDate \
12346.0	United Kingdom	1/18/11 10:01
12347.0	Iceland	12/7/10 14:5712/7/10 14:5712/7/10 14:5712/7/10...
12348.0	Finland	12/16/10 19:0912/16/10 19:0912/16/10 19:0912/1...
12349.0	Italy	11/21/11 9:5111/21/11 9:5111/21/11 9:5111/21/1...
12350.0	Norway	2/2/11 16:012/2/11 16:012/2/11 16:012/2/11 16:...

CustomerID	Country	UnitPrice
12346.0	United Kingdom	1.04
12347.0	Iceland	481.21

12348.0	Finland	178.71
12349.0	Italy	605.10
12350.0	Norway	65.30

Step 7.1.2 Think about what that piece of code does and display the dtype of UnitPrice

```
[ ]: customers.UnitPrice.dtype
```

```
dtype('float64')
```

Step 7.1.3 Pull data from online_rt for CustomerIDs 12346.0 and 12347.0.

```
[ ]: display(online_rt[online_rt.CustomerID == 12347.0].
            sort_values(by='UnitPrice', ascending = False).head())
display(online_rt[online_rt.CustomerID == 12346.0].
            sort_values(by='UnitPrice', ascending = False).head())
```

	InvoiceNo	StockCode	Description	Quantity	\
428966	573511	22423	REGENCY CAKESTAND 3 TIER	6	
286637	562032	22423	REGENCY CAKESTAND 3 TIER	3	
72267	542237	22423	REGENCY CAKESTAND 3 TIER	3	
148300	549222	22423	REGENCY CAKESTAND 3 TIER	3	
428967	573511	23173	REGENCY TEAPOT ROSES	2	

	InvoiceDate	UnitPrice	CustomerID	Country
428966	10/31/11 12:25	12.75	12347.0	Iceland
286637	8/2/11 8:48	12.75	12347.0	Iceland
72267	1/26/11 14:30	12.75	12347.0	Iceland
148300	4/7/11 10:43	12.75	12347.0	Iceland
428967	10/31/11 12:25	9.95	12347.0	Iceland

	InvoiceNo	StockCode	Description	Quantity	\
61619	541431	23166	MEDIUM CERAMIC TOP STORAGE JAR	74215	

	InvoiceDate	UnitPrice	CustomerID	Country
61619	1/18/11 10:01	1.04	12346.0	United Kingdom

Step 7.2 Reinterpreting the initial problem. To reiterate the question that we were dealing with:

“Create a scatterplot with the Quantity per UnitPrice by CustomerID for the top 3 Countries”

The question is open to a set of different interpretations. We need to disambiguate.

We could do a single plot by looking at all the data from the top 3 countries. Or we could do one plot per country. To keep things consistent with the rest of the exercise, let’s stick to the latter option. So that’s settled.

But “top 3 countries” with respect to what? Two answers suggest themselves: Total sales volume (i.e. total quantity sold) or total sales (i.e. revenue). This exercise goes for sales volume, so let’s stick to that.

Step 7.2.1 Find out the top 3 countries in terms of sales volume.

```
[ ]: sales_volume = online_rt.groupby('Country').Quantity.sum().
      ↪sort_values(ascending=False)

top3 = sales_volume.index[1:4]
top3
```

```
Index(['Netherlands', 'EIRE', 'Germany'], dtype='object', name='Country')
```

Step 7.2.2 Now that we have the top 3 countries, we can focus on the rest of the problem: “Quantity per UnitPrice by CustomerID”.

We need to unpack that.

“by CustomerID” part is easy. That means we’re going to be plotting one dot per CustomerID’s on our plot. In other words, we’re going to be grouping by CustomerID.

“Quantity per UnitPrice” is trickier. Here’s what we know:

One axis will represent a Quantity assigned to a given customer. This is easy; we can just plot the total Quantity for each customer.

The other axis will represent a UnitPrice assigned to a given customer. Remember a single customer can have any number of orders with different prices, so summing up prices isn’t quite helpful. Besides it’s not quite clear what we mean when we say “unit price per customer”; it sounds like price of the customer! A reasonable alternative is that we assign each customer the average amount each has paid per item. So let’s settle that question in that manner.

Step 7.3 Modify, select and plot data

Step 7.3.1 Add a column to `online_rt` called **Revenue** calculate the revenue (Quantity * UnitPrice) from each sale. We will use this later to figure out an average price per customer.

```
[ ]: online_rt['Revenue'] = online_rt.Quantity * online_rt.UnitPrice
      online_rt.head()
```

	InvoiceNo	StockCode	Description	Quantity	\
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER	6	
1	536365	71053	WHITE METAL LANTERN	6	
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER	8	
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE	6	
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.	6	

	InvoiceDate	UnitPrice	CustomerID	Country	Revenue
0	12/1/10 8:26	2.55	17850.0	United Kingdom	15.30
1	12/1/10 8:26	3.39	17850.0	United Kingdom	20.34
2	12/1/10 8:26	2.75	17850.0	United Kingdom	22.00
3	12/1/10 8:26	3.39	17850.0	United Kingdom	20.34
4	12/1/10 8:26	3.39	17850.0	United Kingdom	20.34

Step 7.3.2 Group by CustomerID and Country and find out the average price (AvgPrice) each customer spends per unit.

```
[ ]: grouped = online_rt[online_rt.Country.isin(top3)].
      ↳groupby(['CustomerID', 'Country'])

plottable = grouped[['Quantity', 'Revenue']].agg('sum')
plottable['AvgPrice'] = plottable.Revenue / plottable.Quantity

plottable['Country', 'Revenue', :] = plottable.index.get_level_values(1)
plottable.head()
```


CustomerID	Country	Quantity	Revenue	AvgPrice (Country, Revenue)
12426.0	Germany	258	582.73	2.258643
12427.0	Germany	533	825.80	1.549343
12468.0	Germany	366	729.54	1.993279
12471.0	Germany	8212	19824.05	2.414034
12472.0	Germany	4148	6572.11	1.584405

Step 7.3.3 Plot

```
[ ]: g = sns.FacetGrid(plottable)
      g.map(plt.scatter, "Quantity", "AvgPrice", alpha=1)

      g.add_legend()
```

<seaborn.axisgrid.FacetGrid at 0x1af1317bec0>



Step 7.4 What to do now? We aren't much better-off than what we started with. The data are still extremely scattered around and don't seem quite informative.

But we shouldn't despair! There are two things to realize: 1) The data seem to be skewed towards the axes (e.g. we don't have any values where Quantity = 50000 and AvgPrice = 5). So that might suggest a trend. 2) We have more data! We've only been looking at the data from 3 different countries and they are plotted on different graphs.


So: we should plot the data regardless of Country and hopefully see a less scattered graph.

Step 7.4.1 Plot the data for each CustomerID on a single graph

```
[ ]: grouped = online_rt.groupby(['CustomerID'])
     plottable = grouped[['Quantity', 'Revenue']].agg('sum')
     plottable['AvgPrice'] = plottable.Revenue / plottable.Quantity

     plt.scatter(plottable.Quantity, plottable.AvgPrice)
     plt.plot()
```

[]



Step 7.4.2 Zoom in so we can see that curve more clearly


```
[ ]: grouped = online_rt.groupby(['CustomerID', 'Country'])
     plottable = grouped.agg({'Quantity': 'sum',
                             'Revenue': 'sum'})
     plottable['AvgPrice'] = plottable.Revenue / plottable.Quantity
```

```
plt.scatter(plottable.Quantity, plottable.AvgPrice)

plt.xlim(-40,2000)
plt.ylim(-1,80)

plt.plot()
```

[]



8.0.9 8. Plot a line chart showing revenue (y) per UnitPrice (x).

Did Step 7 give us any insights about the data? Sure! As average price increases, the quantity ordered decreases. But that's hardly surprising. It would be surprising if that wasn't the case!

Nevertheless the rate of drop in quantity is so drastic, it makes me wonder how our revenue changes with respect to item price. It would not be that surprising if it didn't change that much. But it would be interesting to know whether most of our revenue comes from expensive or inexpensive items, and how that relation looks like.

That is what we are going to do now.

8.1 Group UnitPrice by intervals of 1 for prices [0,50), and sum Quantity and Revenue.

```
[ ]: price_start = 0
      price_end = 50
      price_interval = 1


      buckets = np.arange(price_start,price_end,price_interval)

      revenue_per_price = online_rt.groupby(pd.cut(online_rt.UnitPrice,
      ↪buckets),observed=False).Revenue.sum()
      revenue_per_price.head()
```

```
UnitPrice
(0, 1]    1107774.544
(1, 2]    2691765.110
(2, 3]    2024143.090
(3, 4]     865101.780
(4, 5]    1219377.050
Name: Revenue, dtype: float64
```

8.3 Plot.

```
[ ]: revenue_per_price.plot()
      plt.xlabel('Unit Price (in intervals of '+str(price_interval)+')')
      plt.ylabel('Revenue')
      plt.show()
```




8.4 Make it look nicer. x-axis needs values.
y-axis isn't that easy to read; show in terms of millions.

```
[ ]: revenue_per_price.plot()

plt.xlabel('Unit Price (in buckets of '+str(price_interval)+'')')
plt.ylabel('Revenue')

plt.xticks(np.arange(price_start,price_end,3),
           np.arange(price_start,price_end,3))
plt.yticks([0, 500000, 1000000, 1500000, 2000000, 2500000],
           ['0', '$0.5M', '$1M', '$1.5M', '$2M', '$2.5M'])
plt.show()
```



[]: